

In [81]:

```
# === 0) Setup
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, StratifiedKFold, cross_validate
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.metrics import (
    roc_auc_score, average_precision_score, confusion_matrix, classification_report, auc,
    precision_recall_curve, roc_curve, accuracy_score, ConfusionMatrixDisplay
)
from sklearn.tree import DecisionTreeClassifier

from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.inspection import permutation_importance

from imblearn.pipeline import Pipeline as ImbPipeline
from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
```

In [82]:

```
#Loading data
df = pd.read_csv('bigml_59c28831336c6604c800002a.csv')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   state            3333 non-null   object 
 1   account length  3333 non-null   int64  
 2   area code        3333 non-null   int64  
 3   phone number    3333 non-null   object 
 4   international plan 3333 non-null   object 
 5   voice mail plan 3333 non-null   object 
 6   number vmail messages 3333 non-null   int64
```

```

7  total day minutes      3333 non-null   float64
8  total day calls        3333 non-null   int64
9  total day charge       3333 non-null   float64
10 total eve minutes      3333 non-null   float64
11 total eve calls        3333 non-null   int64
12 total eve charge       3333 non-null   float64
13 total night minutes    3333 non-null   float64
14 total night calls      3333 non-null   int64
15 total night charge     3333 non-null   float64
16 total intl minutes     3333 non-null   float64
17 total intl calls       3333 non-null   int64
18 total intl charge      3333 non-null   float64
19 customer service calls 3333 non-null   int64
20 churn                  3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

```

```
In [83]: #checking the distributions of the data
df['churn'].value_counts()
```

```
Out[83]: False    2850
True     483
Name: churn, dtype: int64
```

```
In [84]: df.corr()
```

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes	total night calls	total night charge	total intl minutes	total
account length	1.000000	-0.012463	-0.004628	0.006216	0.038470	0.006214	-0.006757	0.019260	-0.006745	-0.008955	-0.013176	-0.008960	0.009514	0.021
area code	-0.012463	1.000000	-0.001994	-0.008264	-0.009646	-0.008264	0.003580	-0.011886	0.003607	-0.005825	0.016522	-0.005845	-0.018288	-0.021
number vmail messages	-0.004628	-0.001994	1.000000	0.000778	-0.009548	0.000776	0.017562	-0.005864	0.017578	0.007681	0.007123	0.007663	0.002856	0.011
total day minutes	0.006216	-0.008264	0.000778	1.000000	0.006750	1.000000	0.007043	0.015769	0.007029	0.004323	0.022972	0.004300	-0.010155	0.001
total day calls	0.038470	-0.009646	-0.009548	0.006750	1.000000	0.006753	-0.021451	0.006462	-0.021449	0.022938	-0.019557	0.022927	0.021565	0.001

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes	total night calls	total night charge	total intl minutes	total
total day charge	0.006214	-0.008264	0.000776	1.000000	0.006753	1.000000	0.007050	0.015769	0.007036	0.004324	0.022972	0.004301	-0.010157	0.000000
total eve minutes	-0.006757	0.003580	0.017562	0.007043	-0.021451	0.007050	1.000000	-0.011430	1.000000	-0.012584	0.007586	-0.012593	-0.011035	0.000000
total eve calls	0.019260	-0.011886	-0.005864	0.015769	0.006462	0.015769	-0.011430	1.000000	-0.011423	-0.002093	0.007710	-0.002056	0.008703	0.010000
total eve charge	-0.006745	0.003607	0.017578	0.007029	-0.021449	0.007036	1.000000	-0.011423	1.000000	-0.012592	0.007596	-0.012601	-0.011043	0.000000
total night minutes	-0.008955	-0.005825	0.007681	0.004323	0.022938	0.004324	-0.012584	-0.002093	-0.012592	1.000000	0.011204	0.999999	-0.015207	-0.010000
total night calls	-0.013176	0.016522	0.007123	0.022972	-0.019557	0.022972	0.007586	0.007710	0.007596	0.011204	1.000000	0.011188	-0.013605	0.000000
total night charge	-0.008960	-0.005845	0.007663	0.004300	0.022927	0.004301	-0.012593	-0.002056	-0.012601	0.999999	0.011188	1.000000	-0.015214	-0.010000
total intl minutes	0.009514	-0.018288	0.002856	-0.010155	0.021565	-0.010157	-0.011035	0.008703	-0.011043	-0.015207	-0.013605	-0.015214	1.000000	0.030000
total intl calls	0.020661	-0.024179	0.013957	0.008033	0.004574	0.008032	0.002541	0.017434	0.002541	-0.012353	0.000305	-0.012329	0.032304	1.000000
total intl charge	0.009546	-0.018395	0.002884	-0.010092	0.021666	-0.010094	-0.011067	0.008674	-0.011074	-0.015180	-0.013630	-0.015186	0.999993	0.030000
customer service calls	-0.003796	0.027572	-0.013263	-0.013423	-0.018942	-0.013427	-0.012985	0.002423	-0.012987	-0.009288	-0.012802	-0.009277	-0.009640	-0.010000
churn	0.016541	0.006174	-0.089728	0.205151	0.018459	0.205151	0.092796	0.009233	0.092786	0.035493	0.006141	0.035496	0.068239	-0.050000

In [85]: #dropping the columns that are not useful for the model

```
columns_to_drop = [
    "phone number",
    "state",
    "area code",
    "total day charge",
    "total eve charge",
    "total night charge",
    "total intl charge"
]

df = df.drop(columns=columns_to_drop)
```

In [86]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   account length   3333 non-null   int64  
 1   international plan 3333 non-null   object  
 2   voice mail plan   3333 non-null   object  
 3   number vmail messages 3333 non-null   int64  
 4   total day minutes  3333 non-null   float64 
 5   total day calls    3333 non-null   int64  
 6   total eve minutes  3333 non-null   float64 
 7   total eve calls    3333 non-null   int64  
 8   total night minutes 3333 non-null   float64 
 9   total night calls   3333 non-null   int64  
 10  total intl minutes 3333 non-null   float64 
 11  total intl calls   3333 non-null   int64  
 12  customer service calls 3333 non-null   int64  
 13  churn              3333 non-null   bool  
dtypes: bool(1), float64(4), int64(7), object(2)
memory usage: 341.9+ KB
```

In [87]: `# Converting categorical variables into numerical values`
`df["international plan"] = df["international plan"].map({"yes":1, "no":0}) # yes stands for churn and No for no churn`
`df["voice mail plan"] = df["voice mail plan"].map({"yes":1, "no":0}) # yes stands for churn and No for no churn`

In [88]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	account length	3333	non-null int64
1	international plan	3333	non-null int64
2	voice mail plan	3333	non-null int64
3	number vmail messages	3333	non-null int64
4	total day minutes	3333	non-null float64
5	total day calls	3333	non-null int64
6	total eve minutes	3333	non-null float64
7	total eve calls	3333	non-null int64
8	total night minutes	3333	non-null float64
9	total night calls	3333	non-null int64
10	total intl minutes	3333	non-null float64
11	total intl calls	3333	non-null int64
12	customer service calls	3333	non-null int64
13	churn	3333	non-null bool

dtypes: bool(1), float64(4), int64(9)
memory usage: 341.9 KB

```
In [89]: df['churn'] = df['churn'].astype(int) # 1 stands for churn and 0 for no churn
```

```
In [90]: # Selecting features and target variable
```

```
X = df.drop(columns=['churn'])
y = df['churn']
```

```
X
```

	account length	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total eve minutes	total eve calls	total night minutes	total night calls	total intl minutes	total intl calls	customer service calls
0	128	0	1	25	265.1	110	197.4	99	244.7	91	10.0	3	1
1	107	0	1	26	161.6	123	195.5	103	254.4	103	13.7	3	1
2	137	0	0	0	243.4	114	121.2	110	162.6	104	12.2	5	0
3	84	1	0	0	299.4	71	61.9	88	196.9	89	6.6	7	2
4	75	1	0	0	166.7	113	148.3	122	186.9	121	10.1	3	3
...
3328	192	0	1	36	156.2	77	215.5	126	279.1	83	9.9	6	2
3329	68	0	0	0	231.1	57	153.4	55	191.3	123	9.6	4	3

	account length	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total eve minutes	total eve calls	total night minutes	total night calls	total intl minutes	total intl calls	customer service calls
3330	28	0	0	0	180.8	109	288.8	58	191.9	91	14.1	6	2
3331	184	1	0	0	213.8	105	159.6	84	139.2	137	5.0	10	2
3332	74	0	1	25	234.4	113	265.9	82	241.4	77	13.7	4	0

3333 rows × 13 columns

In [91]: `# splitting the data into training and testing sets`

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

`# Checking the info and size of the training and testing sets`

```
X_test.info()
```

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 667 entries, 601 to 1962
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   account length    667 non-null    int64  
 1   international plan 667 non-null    int64  
 2   voice mail plan    667 non-null    int64  
 3   number vmail messages 667 non-null    int64  
 4   total day minutes   667 non-null    float64 
 5   total day calls     667 non-null    int64  
 6   total eve minutes   667 non-null    float64 
 7   total eve calls     667 non-null    int64  
 8   total night minutes 667 non-null    float64 
 9   total night calls   667 non-null    int64  
 10  total intl minutes  667 non-null    float64 
 11  total intl calls   667 non-null    int64  
 12  customer service calls 667 non-null    int64  
dtypes: float64(4), int64(9)
memory usage: 73.0 KB
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2666 entries, 3286 to 2762
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   account length    2666 non-null    int64
```

```
1 international plan      2666 non-null  int64
2 voice mail plan        2666 non-null  int64
3 number vmail messages  2666 non-null  int64
4 total day minutes     2666 non-null  float64
5 total day calls        2666 non-null  int64
6 total eve minutes      2666 non-null  float64
7 total eve calls        2666 non-null  int64
8 total night minutes    2666 non-null  float64
9 total night calls      2666 non-null  int64
10 total intl minutes     2666 non-null  float64
11 total intl calls       2666 non-null  int64
12 customer service calls 2666 non-null  int64
dtypes: float64(4), int64(9)
memory usage: 291.6 KB
```

In [92]: `X.describe()`

Out[92]:

	account length	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total eve minutes	total eve calls	total night minutes	total night calls	total intl minutes
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	0.096910	0.276628	8.099010	179.775098	100.435644	200.980348	100.114311	200.872037	100.107711	10.237294
std	39.822106	0.295879	0.447398	13.688365	54.467389	20.069084	50.713844	19.922625	50.573847	19.568609	2.791840
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	23.200000	33.000000	0.000000
25%	74.000000	0.000000	0.000000	0.000000	143.700000	87.000000	166.600000	87.000000	167.000000	87.000000	8.500000
50%	101.000000	0.000000	0.000000	0.000000	179.400000	101.000000	201.400000	100.000000	201.200000	100.000000	10.300000
75%	127.000000	0.000000	1.000000	20.000000	216.400000	114.000000	235.300000	114.000000	235.300000	113.000000	12.100000
max	243.000000	1.000000	1.000000	51.000000	350.800000	165.000000	363.700000	170.000000	395.000000	175.000000	20.000000

In [93]: `# We need to standardize the data before feeding it to the model`
`scaler = StandardScaler()`
`X_train_scaled = scaler.fit_transform(X_train)`
`X_test_scaled = scaler.transform(X_test)`

In [94]: `#fit the model and balance the highly imbalanced classes using class_weight parameter`
`model = LogisticRegression(class_weight="balanced",random_state=42)`

```
model.fit(X_train_scaled, y_train)
```

```
Out[94]: LogisticRegression(class_weight='balanced', random_state=42)
```

```
In [95]: #make predictions
```

```
y_pred = model.predict(X_test_scaled)
y_pred_proba = model.predict_proba(X_test_scaled)[:, 1]
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.76	0.84	570
1	0.35	0.74	0.47	97
accuracy			0.76	667
macro avg	0.65	0.75	0.66	667
weighted avg	0.86	0.76	0.79	667

```
In [96]: # Evaluate the model using ROC AUC and Average Precision Score
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
print("ROC AUC:", roc_auc_score(y_test, y_pred_proba))
```

Accuracy: 0.7586206896551724

	precision	recall	f1-score	support
0	0.95	0.76	0.84	570
1	0.35	0.74	0.47	97
accuracy			0.76	667
macro avg	0.65	0.75	0.66	667
weighted avg	0.86	0.76	0.79	667

ROC AUC: 0.8151021884608429

Interpretaion:

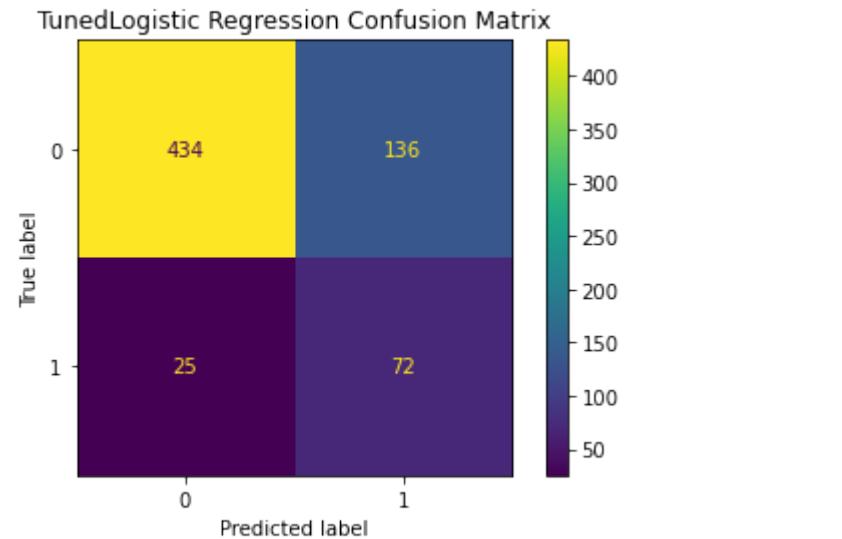
The model shows good discrimination ability (ROC AUC 0.815) but struggles with precision for the minority class (class 1). This means it identifies positives well but also produces many false positives.

```
In [97]: cm = confusion_matrix(y_test, y_pred)
```

```
fig, ax = plt.subplots()
```

```
display = ConfusionMatrixDisplay(confusion_matrix=cm)
display.plot(ax=ax)

ax.set_title("Tuned Logistic Regression Confusion Matrix")
plt.show()
```



In [98]:

```
TP = 72
FP = 136
TN = 434
FN = 25
```

```
Precision = TP / (TP + FP)
Recall = TP / (TP + FN)
F1_Score = 2 * (Precision * Recall) / (Precision + Recall)
print(f"Precision: {Precision:.3f}")
print(f"Recall: {Recall:.3f}")
print(f"F1 Score: {F1_Score:.3f}")
```

```
Precision: 0.346
Recall: 0.742
F1 Score: 0.472
```

Interpretation:

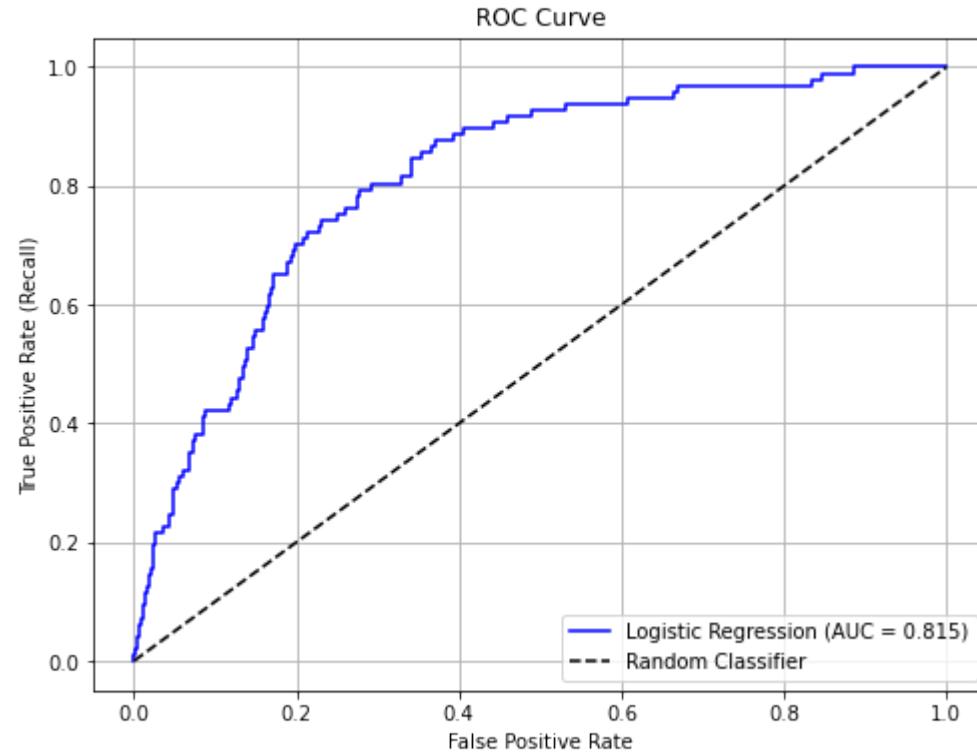
F1 score is low since the precision score is low and the score of 0.472 , while recall is at 0.742 meaning it correctly identifies 74% of customers who actual churned. Having a low precison means it picks more clients as churned thus the company using its resources to deal with clients who were not churning.

```
In [99]: roc_auc = roc_auc_score(y_test, y_pred_proba)
print("ROC AUC Score:", roc_auc)

fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)

plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'Logistic Regression (AUC = {roc_auc:.3f})', color='blue')
plt.plot([0,1], [0,1], 'k--', label='Random Classifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate (Recall)')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```

ROC AUC Score: 0.8151021884608429



Decision tree

In [100...]

```
#instantiate our tree and balancing the imbalanced data
dt = DecisionTreeClassifier(
    max_depth=5,
    min_samples_split=20,
    class_weight="balanced",
    random_state=42
) # here we are still balancing the data as the data for target has class heavily imbalanced

dt.fit(X_train, y_train)
```

Out[100...]

```
DecisionTreeClassifier(class_weight='balanced', max_depth=5,
                      min_samples_split=20, random_state=42)
```

In [101...]

```
#Predict on the test data
y_pred_dt = dt.predict(X_test)
y_prob_dt = dt.predict_proba(X_test)[:, 1]
```

```
#Evaluate the model
print("Classification Report:\n", classification_report(y_test, y_pred_dt))
print("Decision Tree Performance:", accuracy_score(y_test, y_pred_dt))
```

```
Classification Report:
precision    recall    f1-score   support
          0       0.95      0.93      0.94      570
          1       0.65      0.73      0.69       97

   accuracy                           0.90      667
  macro avg       0.80      0.83      0.82      667
weighted avg       0.91      0.90      0.91      667
```

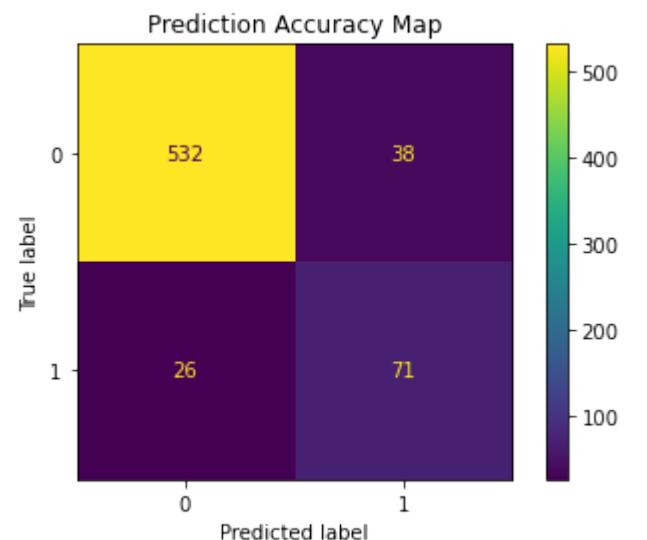
```
Decision Tree Performance: 0.904047976011994
```

```
In [108...]: cm = confusion_matrix(y_test, y_pred_dt)

fig, ax = plt.subplots()

display = ConfusionMatrixDisplay(confusion_matrix=cm)
display.plot(ax=ax)

ax.set_title("Prediction Accuracy Map")
plt.show()
```



Interpretation

The weighting did improve recall for churners (0.73), meaning the model is catching more of the clients at risk of leaving.

Precision for churn (0.65) is still relatively low, so the model raises some false alarms — but in churn prediction, false positives are usually less costly than false negatives (better to flag a client who won't leave than miss one who will).

The model is now more balanced: it doesn't just favor the majority class (non-churn), but still struggles to achieve the same precision/recall levels for churn as for non-churn.

In [103...]

```
# Compare both models using ROC

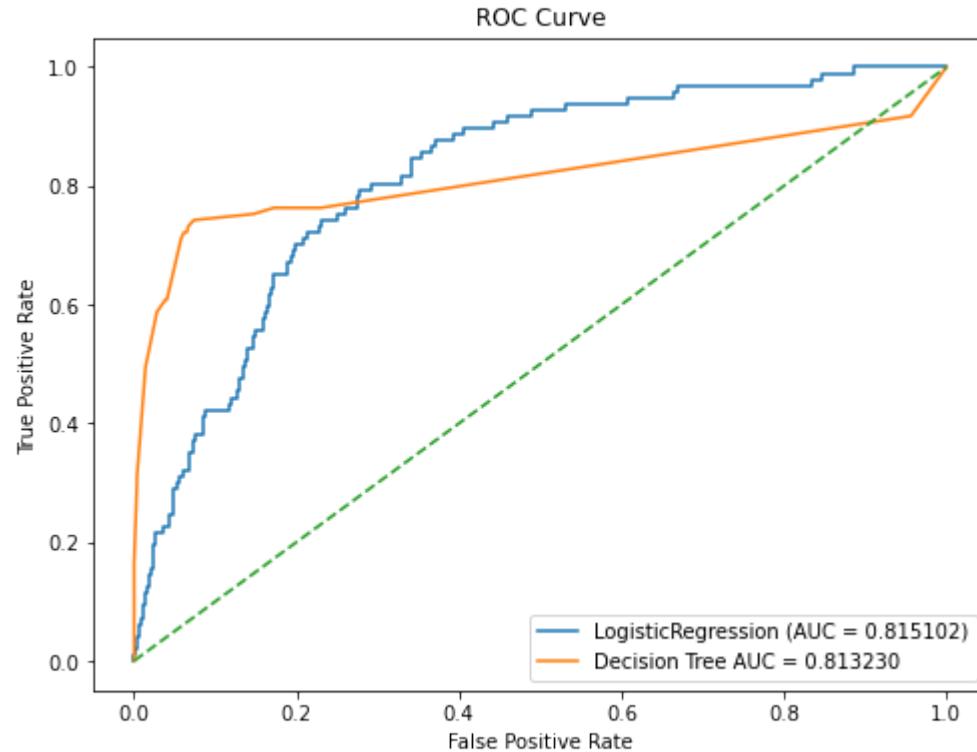
lr_fpr, lr_tpr, i = roc_curve(y_test, y_pred_proba)
dt_fpr, dt_tpr, i = roc_curve(y_test, y_prob_dt)

lr_auc = auc(lr_fpr, lr_tpr)
dt_auc = auc(dt_fpr, dt_tpr)
```

In [107...]

```
plt.figure(figsize=(8, 6))
plt.plot(lr_fpr, lr_tpr, label=f'LogisticRegression (AUC = {lr_auc:2f})')
plt.plot(dt_fpr, dt_tpr, label=f'Decision Tree AUC = {dt_auc:2f}')
plt.plot([0,1], [0, 1], linestyle='--')

plt.xlabel('False Positive Rate')
plt.title('ROC Curve')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



Conclusion

Both the logistic regression and the decision tree model achieved an AUC of 0.81, indicating good discrimination ability. The models can correctly distinguish between a churned client and a non-churned client, 81% of the time thus they can be used to support customer retention strategy.